

# A very introduction to dplyr

Rodrigo Esteves de Lima-Lopes  
State University of Campinas  
[rll307@unicamp.br](mailto:rll307@unicamp.br)

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Using dplyr</b>	<b>2</b>
2.1	Loading the package . . . . .	2
2.2	Loading data . . . . .	2
2.3	Filter . . . . .	2
2.4	Select . . . . .	4
2.5	Rename . . . . .	4
2.6	Mutate . . . . .	4
2.7	Arrange and filter . . . . .	6

## 1 Introduction

You might have observed that some commands showed up in our last tutorial:

1. `select()`
2. `mutate()`
3. `rename()`
4. `count()`
5. `arrange()`
6. `filter()`

Each of these commands (and many others, there is an cheatsheet for you) is loaded as we execute `library(dplyr)`. `dplyr` is part of the [tidyverse](#) collection of packages for data science. They make the task to manipulate data easier than if we were using ordinary R commands. `dplyr` is usually referred as ‘a grammar of data manipulation’.

However, please, keep in mind that every task in R might be performed in different ways, using different commands or packages.

We will discuss here some of the packages utilities, not all. If you need further information, please refer to the [dplyr website](#) or to the readme at [CRAN](#). Remember that after loading the package, you can always type `?+command` to get further information.

## 2 Using dplyr

### 2.1 Loading the package

We load the package using the following command:

```
library(dplyr)
```

### 2.2 Loading data

During this tutorial we are going to use `starwars`, a data frame with information regarding the film series (you can tell your professor is a geek). This data frame comes when you load the packages, and it is common that some packages bring some data for training purposes. So, run:

```
head(starwars,10)
```

```
## # A tibble: 10 x 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>    <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Luke S~    172    77 blond      fair        blue        19    male masculi
## 2 C-3PO     167    75 <NA>       gold        yellow      112    none masculi
## 3 R2-D2      96    32 <NA>       white, bl~ red         33    none masculi
## 4 Darth ~   202   136 none       white       yellow     41.9    male masculi
## 5 Leia O~   150    49 brown      light       brown       19    fema~ femini
## 6 Owen L~   178   120 brown, grey light       blue        52    male masculi
## 7 Beru W~   165    75 brown      light       blue        47    fema~ femini
## 8 R5-D4      97    32 <NA>       white, red red         NA     none masculi
## 9 Biggs ~   183    84 black      light       brown       24    male masculi
## 10 Obi-Wa~  182    77 auburn, wh~ fair        blue-gray   57    male masculi
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

The data frame brings some basic information about some of the main characters at the series. Here we printed only the 10 first rows. Let us run a `summary()` and a `str()` to check its structure:

```
str(starwars)
```

```
summary(starwars)
```

We have pretty long results, but some conclusions might be interesting:

- It is a data frame
- It has numeric, list and character columns
- It has only the seven main movies.

### 2.3 Filter

**Filter** is a command that might be used for filtering data based on a data frame. For example, I will filter all the “droids” in the classical movies:

```
head(starwars %>%
  filter(species == "Droid"))
```

```
## # A tibble: 6 x 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>    <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 C-3PO      167    75 <NA>      gold        yellow        112 none masculi~
## 2 R2-D2       96    32 <NA>      white, blue red          33 none masculi~
## 3 R5-D4       97    32 <NA>      white, red  red          NA none masculi~
## 4 IG-88      200   140 none      metal       red          15 none masculi~
## 5 R4-P17      96    NA none      silver, red red, blue    NA none feminine
## 6 BB8        NA    NA none      none        black        NA none masculi~
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

*head()* and *tail()* are important commands in **R** language, they will show us the top and bottom **6** lines of any variable I print on my terminal. There is also a second argument, numeric, that may increase or decrease the number of lines. Now let us see all humans.

```
head(starwars %>%
  filter(species == "Human"), 3)
```

```
## # A tibble: 3 x 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>    <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Luke Sk~    172    77 blond     fair        blue        19  male masculi~
## 2 Darth V~    202   136 none      white       yellow       41.9 male masculi~
## 3 Leia Or~    150    49 brown     light       brown        19  female femini~
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

Now I have the first three columns. Mind you that the number *3* came in a specific place inside the parenthesis: `command1(command2(3))`. It is so because it is part of `command1()`. The position of the arguments is something we have to keep in mind when we embed commands.

Any logical operator can combine into a search:

```
head(starwars %>%
  filter(species == "Human" & hair_color=="none"), 3)
```

```
## # A tibble: 3 x 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>    <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Darth V~    202   136 none      white       yellow       41.9 male masculi~
## 2 Lobot       175    79 none      light       blue         37  male masculi~
## 3 Mace Wi~    188    84 none      dark        brown        72  male masculi~
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

## 2.4 Select

Select help me to get just a couple of information and display it

```
head(starwars %>%
  select(name, ends_with("color")))
```

```
## # A tibble: 6 x 4
##   name      hair_color skin_color eye_color
##   <chr>      <chr>      <chr>      <chr>
## 1 Luke Skywalker blond      fair      blue
## 2 C-3PO      <NA>      gold      yellow
## 3 R2-D2      <NA>      white, blue red
## 4 Darth Vader none      white      yellow
## 5 Leia Organa brown      light      brown
## 6 Owen Lars  brown, grey light      blue
```

## 2.5 Rename

Rename help us to change easily the name of a single column

```
head(rename(flights, airline_car = carrier))
```

```
## # A tibble: 6 x 19
##   year month day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1   1     517             515           2     830           819
## 2 2013     1   1     533             529           4     850           830
## 3 2013     1   1     542             540           2     923           850
## 4 2013     1   1     544             545          -1    1004          1022
## 5 2013     1   1     554             600          -6     812           837
## 6 2013     1   1     554             558          -4     740           728
## # ... with 11 more variables: arr_delay <dbl>, airline_car <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

## 2.6 Mutate

`mutate()` allows me to create new columns based on some criteria I need. It draws data from existing columns and executes commands to add a new data column in my data frame. For this command let us use another built in R data frame called `nycflights13`:

```
library(nycflights13)
summary(flights)
```

```
##      year      month      day      dep_time      sched_dep_time
## Min.   :2013   Min.   : 1.000   Min.   : 1.00   Min.   : 1      Min.   : 106
## 1st Qu.:2013   1st Qu.: 4.000   1st Qu.: 8.00   1st Qu.: 907    1st Qu.: 906
## Median :2013   Median : 7.000   Median :16.00   Median :1401    Median :1359
## Mean   :2013   Mean   : 6.549   Mean   :15.71   Mean   :1349    Mean   :1344
```

```
## 3rd Qu.:2013 3rd Qu.:10.000 3rd Qu.:23.00 3rd Qu.:1744 3rd Qu.:1729
## Max. :2013 Max. :12.000 Max. :31.00 Max. :2400 Max. :2359
## NA's :8255
## dep_delay arr_time sched_arr_time arr_delay
## Min. : -43.00 Min. : 1 Min. : 1 Min. : -86.000
## 1st Qu.: -5.00 1st Qu.:1104 1st Qu.:1124 1st Qu.: -17.000
## Median : -2.00 Median :1535 Median :1556 Median : -5.000
## Mean : 12.64 Mean :1502 Mean :1536 Mean : 6.895
## 3rd Qu.: 11.00 3rd Qu.:1940 3rd Qu.:1945 3rd Qu.: 14.000
## Max. :1301.00 Max. :2400 Max. :2359 Max. :1272.000
## NA's :8255 NA's :8713 NA's :9430
## carrier flight tailnum origin
## Length:336776 Min. : 1 Length:336776 Length:336776
## Class :character 1st Qu.: 553 Class :character Class :character
## Mode :character Median :1496 Mode :character Mode :character
## Mean :1972
## 3rd Qu.:3465
## Max. :8500
## dest air_time distance hour
## Length:336776 Min. : 20.0 Min. : 17 Min. : 1.00
## Class :character 1st Qu.: 82.0 1st Qu.: 502 1st Qu.: 9.00
## Mode :character Median :129.0 Median : 872 Median :13.00
## Mean :150.7 Mean :1040 Mean :13.18
## 3rd Qu.:192.0 3rd Qu.:1389 3rd Qu.:17.00
## Max. :695.0 Max. :4983 Max. :23.00
## NA's :9430
## minute time_hour
## Min. : 0.00 Min. :2013-01-01 05:00:00
## 1st Qu.: 8.00 1st Qu.:2013-04-04 13:00:00
## Median :29.00 Median :2013-07-03 10:00:00
## Mean :26.23 Mean :2013-07-03 05:22:54
## 3rd Qu.:44.00 3rd Qu.:2013-10-01 07:00:00
## Max. :59.00 Max. :2013-12-31 23:00:00
##
```

If you do not have it installed, please do it. The instructions are in a previous tutorial. `nycflights13` brings all data on Flights in New York airports during the year of 2013, and it is a great tool for training data manipulation. The data brings some information regarding the delays (arrivals and departures). I want a new column that will show me the total of delay a plane might have.

```
mutate(flights, total_delay = arr_delay+dep_delay)%>%
  select(total_delay)%>%
  head()
```

```
## # A tibble: 6 x 1
##   total_delay
##   <dbl>
## 1         13
## 2         24
## 3         35
## 4        -19
## 5        -31
## 6          8
```

Notice that here instead of embedding the command, the *pipe* `%>%` was my syntax choice. Initially the *pipe* `%>%` function was introduced in R by the package `magrittr` which is able to do much much more than just pipening our code.

If I want to select more than a column I will have to save it as a variable:

```
delayed_flights <- mutate(flights, total_delay = arr_delay + dep_delay)
head(select(delayed_flights, carrier, total_delay))
```

```
## # A tibble: 6 x 2
##   carrier total_delay
##   <chr>      <dbl>
## 1 UA          13
## 2 UA          24
## 3 AA          35
## 4 B6         -19
## 5 DL         -31
## 6 UA           8
```

However, from R 4.10 on, it is now possible to use a native *pipe* operator `|>`. The reasons to such change is that the use of `%>%` takes much more in terms of memory than any R native resources. There are two ways of activating native `|>`:

- Get any version over R 4.10 that it will work out of the box
  - For MacOS users: it has been just updated
- Activate on preferences `>` code

If you want to see a complete discussion on the differences between the two approaches, I would recommend the following [video](#)

So, from now on I will be using `%>%` or `|>` interchangeably.

## 2.7 Arrange and filter

To discuss arrange and filter I will get back to our previous Gutenberg project data frame:

```
geral.list <- geral.list.df %>%
  unnest_tokens(word, text) %>%
  count(word, sort = TRUE) %>%
  anti_join(my.stopwords, by= "word") %>%
  mutate((freq = n / sum(n))*100) %>%
  arrange(desc(freq))
colnames(geral.list) <- c('word', 'n', 'freq')
```

- `anti_join`: excludes whatever is in the column `word` in `my.stopwords` file.
- `arrange`: arrange the data according to some criterion, here the column `freq` in descending order.

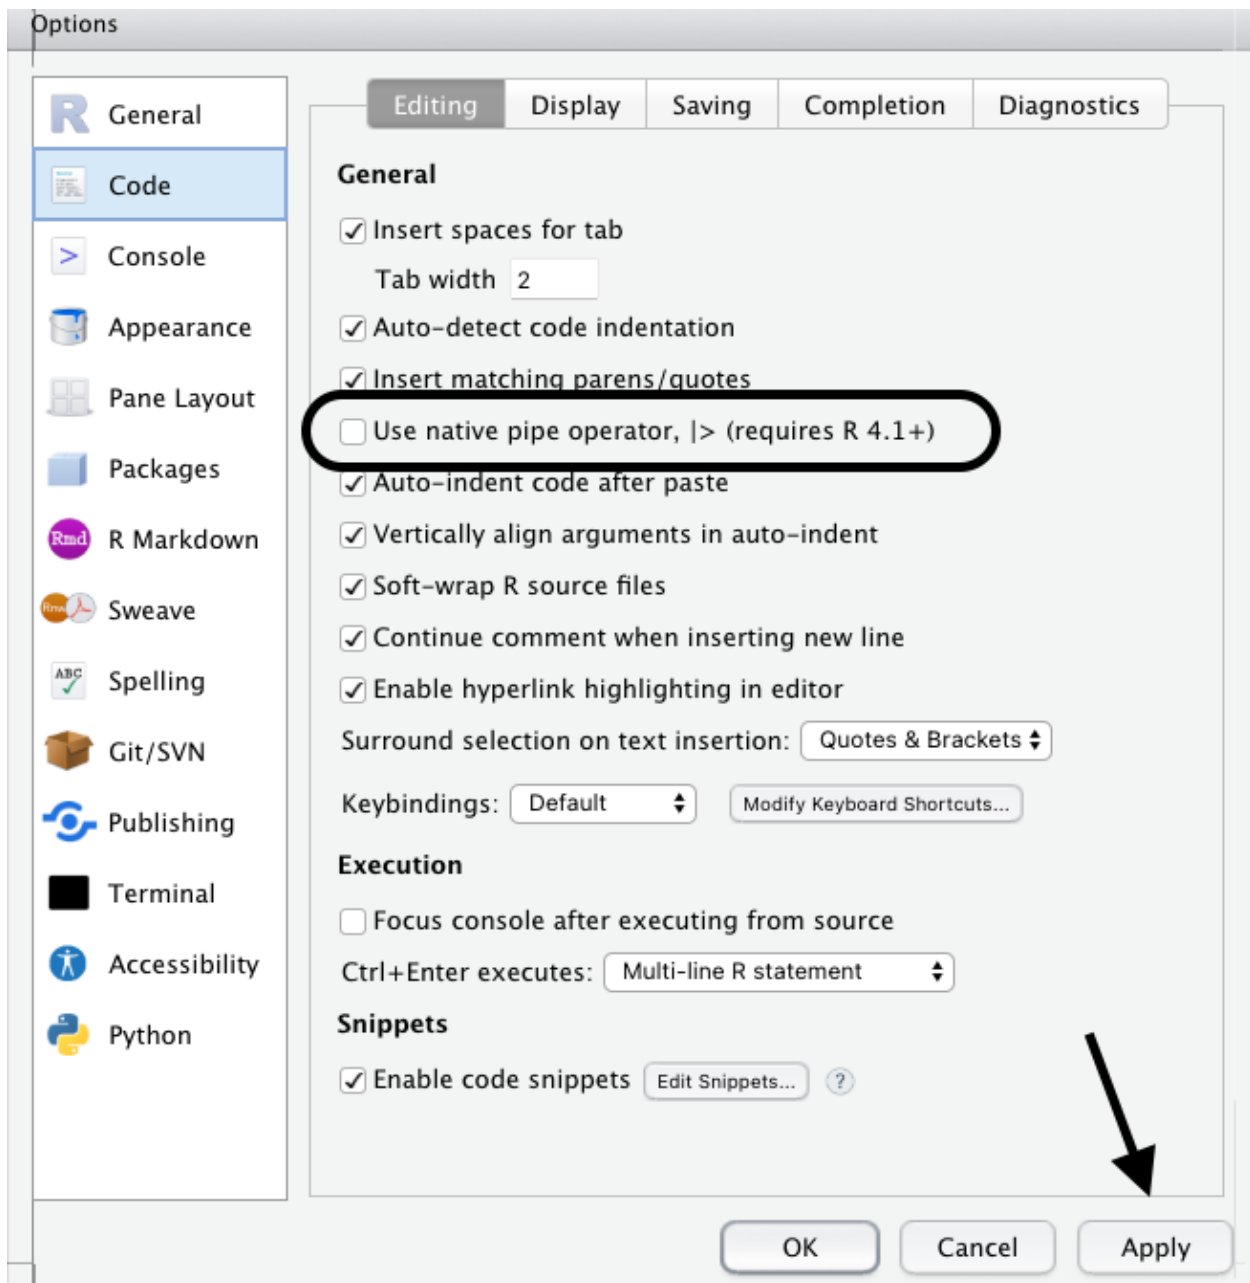


Figure 1: Native Pipe