

# TidyText: Reviewing in a set of Clarisse Lispector Novels

Rodrigo Esteves de Lima-Lopes  
State University of Campinas  
rll307@unicamp.br

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Purpose of this notebook . . . . .	1
1.2	Clarisse Lispector . . . . .	2
<b>2</b>	<b>Loading R packages</b>	<b>2</b>
<b>3</b>	<b>Loading data</b>	<b>2</b>
3.1	Data and stop-words . . . . .	2
3.2	Cleaning data . . . . .	3
<b>4</b>	<b>Analysis</b>	<b>3</b>
4.1	Create and compare a wordlist . . . . .	3
4.1.1	Visualising our wordlist . . . . .	5
4.2	Analysing bigrams . . . . .	5
4.2.1	Going forward . . . . .	7

## 1 Introduction

### 1.1 Purpose of this notebook

In this post I am going to discuss some strategies of comparison between texts. It was produced in order to assist colleagues who work in the area of Corpus Linguistics and Systemic Functional Linguistics, as a way to use R in their research. It is part of my CNPq-funded project and seeks to make corpus tools and network analysis accessible. If you have any doubts or wish to make any research contact please send me an email.

This document is based mostly in the package Tidytext and in the following book:

Silge, Julia and David Robinson. 2017. *Text Mining with R: A Tidy Approach*. First edition. Beijing/Boston: O'Reilly.

## 1.2 Clarisse Lispector

In this notebook I am going to compare two novels by Clarisse Lispector: A Hora da Estrela and A Paixão segundo GH in their original Portuguese version. Some more information about Lispector might be found [here](#).

## 2 Loading R packages

In this article we are going to use the following packages:

```
library(tidytext) #OK
library(dplyr) #OK
library(tidyverse)
library(tidyr)
library(stringr)
# library(ggraph) Only one function in the code
# # library(igraph) Only one function in the code
# library(textreadr) only one function in the code
# library(abjutils) only one command inside functions
# library(textclean) only one command inside functions
# library(readr) only one function in the code
# library(scales) Only one function in the code
# library(ggplot2) Only one function in the code
```

Each of these packages has a specific function:

1. **tidytext**: Text manipulation
2. **dplyr**: Data manipulation, such as tables and data frames
3. **stringr**: Manipulation of strings
4. **tidyverse**: Text manipulation, it loads some functions from other packages
5. **tidyr**: Data manipulation, such as tables and data frames

Since we are going to use just some functions, these packages are not going to be loaded, we are going to call their functions at the right moment:

1. **readr** and **textreadr**: Read documents and data into R
2. **scales**: Data representation
3. **abjutils**: Data manipulation for Brazilian attorneys of law
4. **textclean**: Clean texts
5. **ggraph** and **igraph**: Preparing and plotting network graphs
6. **ggplot2**: Doing some common plots.

## 3 Loading data

### 3.1 Data and stop-words

Since my focus here is on the semantically relevant Lexis, I will load a list of stop words, in order to filter prepositions, articles and conjunctions. There are a number of stop-words list for Portuguese, I would advise you to prepare your own based on your research objectives. The code bellow also loads the two novels in R.

```
My.Stopwords <- data.frame(
  readr::read_csv("stop_port2.csv", col_names = FALSE), stringsAsFactors = FALSE
)
colnames(My.Stopwords) <- "word"
```

```
head(My.Stopwords, 10)
```

Now we are going to load the novels:

```
estrela <- textreadr::read_document("estrela.txt")
```

```
paixao <- textreadr::read_document("paixao.txt")
```

## 3.2 Cleaning data

Our first step is going to clean our textual data. For this study we will:

1. Remove diacritics
  - This is not always necessary, it will depend of our objectives
1. Replace possible HTML codes
2. Change words to lower case
3. Delete numerals
4. Delete extra spaces and extra lines

To clean the data we going to import a previously function I have developed earlier. This function is in another R script, called `05_TidyText_functions.R`. We going to load it using the following command:

```
source("05_TidyText_functions.R")
```

## 4 Analysis

### 4.1 Create and compare a wordlist

Our first step will be processing the two novels in a set of words with their frequencies.

```
estrela <- data.frame(text = estrela, stringsAsFactors = F)
```

```
paixao.tidy <- analyse.paixao |>
  unnest_tokens(word, text) |>
  anti_join(My.Stopwords, by = "word")
```

```
paixao <- data.frame(text = paixao, stringsAsFactors = F)
```

```
paixao.tidy <- paixao |>
  unnest_tokens(word, text) |>
  anti_join(My.Stopwords, by = "word")
```

The code above creates a data frame from each novel. It is important to make the files processable by R. Our next step is to apply our customised function (`Clean_String`) to the files. Note that we keep the data frame format.

```
analise.paixao <- Clean_String(paixao$text) |>
  data.frame()
colnames(analise.paixao) <- 'text'

analise.estrela <- Clean_String(estrela$text) |>
  data.frame()
colnames(analise.estrela) <- 'text'
```

Now our next step is prepare the wordlist. The following lines process the data frame in order to:

- Make it in a word per line forma
- Filter stop-words out
- Count and present the wordlist

```
paixao.tidy <- analise.paixao |>
  unnest_tokens(word, text) |>
  anti_join(My.Stopwords, by = "word")

estrela.tidy <- analise.estrela |>
  unnest_tokens(word, text) |>
  anti_join(My.Stopwords, by = "word")
head(paixao.tidy)
head(estrela.tidy)
```

As we can see, there is a column which displays a sequential number, representing the line of the list, and another which displays the words itself. However it is also possible to observe that *procurando* is present twice in the list. It happens because we are not counting occurrences, but only displaying the novel's words.

```
paixao.l <- paixao.tidy |>
  count(word, sort = TRUE)
```

```
esterla.l <- estrela.tidy |>
  count(word, sort = TRUE)
```

The code above transforms these lists in something similar to a word list. The first line creates the file itself, which I chose to name with the suffix “\*.l”. The operator “|>” is part of the tidyverse, its function is to help us to “pass” one line of code through to the other and perform more than a command at once. The code counts the words and creates a new column with the results.

Our next step is to build a table comparing the frequency of such words. It is necessary in order to generate the visualisations and make the analysis possible. In the first line we create a new object and add two columns in it, one for each book, and make sure that the list of words identify to the book where it is from. I the

following line we clear the text, keeping only alphabetical words and deleting numbers and special symbols. We group the words by book (*livro* in Portuguese) and calculate the relative proportion of the words within each novel. Since absolute values do not mean much, we delete them.

```
frequencia.clarisse <- bind_rows(mutate(paixao.tidy, livro = "P"),
                                mutate(estrela.tidy, livro = "H")) |>
  mutate(word = str_extract(word, "[a-z']+")) |>
  count(livro, word) |>
  group_by(livro) |>
  mutate(proportion = (n / sum(n))*100) |>
  select(-n) |>
  spread(livro, proportion)
head(frequencia.clarisse)
```

It was possible to observe that the word **aaaar** is present in the top of the list. If we want to clean it out, we have to add it to our stop wordlist:

```
My.Stopwords <- data.frame(word = c(My.Stopwords$word, "aaaar"))
```

Then we rerun the four pieces of code above.

#### 4.1.1 Visualising our wordlist

Now it is time to compare the Lexis in the two books. We will do so by using a ggplot command. Mostly we are putting both novels side by side, making them overlap. The position is calculated by a log scale that helps to represent the overlapping. The more central a word is in the graphic, the less “specific” of a book it is.

```
ggplot(frequencia.clarisse, aes(x = H, y = P,
                                color = abs(H - P))) +
  geom_abline(color = "gray40", lty = 2) +
  geom_jitter(alpha = 0.1, size = 2.5, width = 0.3, height = 0.3) +
  geom_text(aes(label = word), check_overlap = TRUE, vjust = 1.5) +
  scale_x_log10(labels = scales::percent_format()) +
  scale_y_log10(labels = scales::percent_format()) +
  scale_color_gradient(limits = c(0, 0.001),
                       low = "darkslategray4", high = "gray75") +
  theme(legend.position = "none") +
  labs(y = "Paixão segundo GH", x = "Hora da Estrela")
```

## 4.2 Analysing bigrams

In order to do a complete bigrams analysis. We will do some changes to the files. First we add a column to each and then make those a single data frame. Note that the book column will become a **factor**

```
# Adding a column with book titles
analise.estrela$book <- "Hora da Estrela"
analise.paixao$book <- "Paixao segundo GH"
# Joining
clarisse.livros <- rbind(analise.estrela, analise.paixao)
clarisse.livros$book <- as.factor(clarisse.livros$book)
head(clarisse.livros)
```

Our next step is to create the bigrams itself.

```
# Creating bigrams
clarisse.bigrams <- clarisse.livros |>
  unnest_tokens(bigram, text, token = "ngrams", n = 2)
bigrams.count <- clarisse.bigrams |>
  count(bigram, sort = TRUE)
head(clarisse.bigrams)
head(bigrams.count)
```

In the first line, the command `unnest_tokens` creates a column which the bigrams found, taking from the column `text`, the `n=2` tells how long my ngram is supposed to be. In the second line, it creates a data frame telling how many of each are. The command `count` does it for me.

Our next step is to separate the words, these will be important for our network analysis. Note that the command is separating two columns using a single space as a criterion.

```
clarisse.separated <- clarisse.bigrams |>
  separate(bigram, c("word1", "word2"), sep = " ")
head(clarisse.separated)
```

Since we will also to analyse each novel separately, we will make the same procedure for each

```
hora.count <- clarisse.bigrams |>
  subset(book == "Hora da Estrela") |>
  count(bigram, sort = TRUE) |>
  separate(bigram, c("word1", "word2"), sep = " ")
hora.count <- hora.count[-3,]

paixao.count <- clarisse.bigrams |>
  subset(book == "Paixao segundo GH") |>
  count(bigram, sort = TRUE) |>
  separate(bigram, c("word1", "word2"), sep = " ")
paixao.count <- paixao.count[-3,]

head(hora.count)
head(paixao.count)
```

Now we are going to clean the stopwords and see the difference

```
clarisse.filtered <- clarisse.separated |>
  filter(!word1 %in% My.Stopwords$word) |>
  filter(!word2 %in% My.Stopwords$word)
head(clarisse.filtered)
```

Now we count the filtered bigrams and have a look at it

```
bigrams.count <- clarisse.filtered |>
  count(word1, word2, sort = TRUE)
head(bigrams.count)
```

### 4.2.1 Going forward

Now we will analyse the relevance of each bigrams in the novel

```
clarisse.united <- clarisse.filtered |>
  unite(bigram, word1, word2, sep = ' ')

clarisse.tf_idf <- clarisse.united |>
  count(book, bigram) |>
  bind_tf_idf(bigram, book, n) |>
  arrange(desc(tf_idf))
head(clarisse.tf_idf)
```

Now we are going to plot these frequencies

```
## Plotting the relative frequency
clarisse.tf_idf |>
  arrange(desc(tf_idf)) |>
  mutate(bigram = factor(bigram, levels = rev(unique(bigram)))) |>
  group_by(book) |>
  top_n(10) |>
  ungroup() |>
  ggplot2::ggplot(aes(bigram, tf_idf, fill = book)) +
  geom_col(show.legend = FALSE) +
  labs(x = NULL, y = "tf-idf") +
  facet_wrap(~book, ncol = 2, scales = "free") +
  coord_flip()
```

Our next step is to prepare and plot a couple of network graphics. First each novel individually:

```
set.seed(1973)
#Selecting the words and creating the interaction file
# Hora da estrela
hora.graph <- hora.count |>
  filter(n > 10) |>
  igraph::graph_from_data_frame()
hora.graph
ggraph::ggraph(hora.graph, layout = "fr") +
  ggraph::geom_edge_link() +
  ggraph::geom_node_point() +
  ggraph::geom_node_text(aes(label = name),
                        vjust = 1, hjust = 1)
```

```
set.seed(1973)
paixao.graph <- paixao.count |>
  filter(n > 20) |>
  igraph::graph_from_data_frame()
print(paixao.graph)
ggraph::ggraph(paixao.graph, layout = "fr") +
  ggraph::geom_edge_link() +
  ggraph::geom_node_point() +
  ggraph::geom_node_text(aes(label = name), vjust = 1, hjust = 1)
```

Now doing both books together

```
set.seed(1973)
# Getting the filtered bigrams
clarisse.g <- clarisse.filtered[,2:3] |>
  count(word1, word2, sort = TRUE)
clarisse.g <- clarisse.g[-1,]
head(clarisse.g)

general.graph <- clarisse.g |>
  filter(n > 4) |>
  igraph::graph_from_data_frame()
print(general.graph)

ggraph::ggraph(general.graph, layout = "fr") +
  ggraph::geom_edge_link() +
  ggraph::geom_node_point() +
  ggraph::geom_node_text(aes(label = name), vjust = 1, hjust = 1)
```